

# The CL-Aff Happiness Shared Task: Results and Key Insights

Kokil Jaidka<sup>1,2</sup>, Saran Mumick<sup>2,3</sup>, Niyati Chhaya<sup>4</sup>, and Lyle Ungar<sup>2</sup>

<sup>1</sup> Nanyang Technological University, Singapore

<sup>2</sup> University of Pennsylvania, USA

<sup>3</sup> Megagon Labs, USA

<sup>4</sup> Adobe Research, India

jaidka@ntu.edu.sg

**Abstract.** This overview describes the official results of the CL-Aff Shared Task 2019 – in Pursuit of Happiness. The dataset comprised a semi-supervised classification task and an open-ended knowledge modeling task on a dataset of over 80,000 brief autobiographical accounts of happy moments, crowdsourced from Amazon Mechanical Turk. The Shared Task was organized as a part of the 2<sup>nd</sup> Workshop on Affective Content Analysis @ AAAAI-19, held in Honolulu, USA on January 27, 2019. This paper compares the participating systems in terms of their accuracy and F-1 scores at predicting two facets of happiness. The complete annotated dataset is available on Harvard Dataverse at <https://goo.gl/3rcZqf>. The annotation instructions and the scripts used for evaluation are available at the Git repository at <https://github.com/kj2013/claff-happydb>.

## 1 Introduction

The purpose of the CL-Aff Shared Task is to challenge the current understanding of emotion through a task that models the experiential, contextual and agentic attributes of happy moments. It has long been known that human affect is context-driven, and that labeled datasets should account for these factors in generating predictive models of affect. The Shared Task is organized in collaboration with researchers at Megagon Labs and builds upon the HappyDB dataset [1], comprising human accounts of ‘happy moments’. The Shared Task comprised of two sub-tasks for analyzing happiness and well-being in written language, on a corpus of over 80,000 descriptions of happy moments, as described here: Given: An account of a happy moment, marked with individual’s demographics, recollection time and relevant labels.

- Task 1: Semi-Supervised classification task - Predict thematic labels (Agency/Sociality) on unseen data, based on a small labeled and large unlabeled training data.<sup>5</sup>

<sup>5</sup> In the annotation task and the Shared Task, the label names we provided were ‘Agency’ and ‘Social’. We have since renamed ‘Social’ to ‘Sociality’ so that both Agency and Sociality can be grammatically consistent.

- Task 2: Suggest interesting ways to automatically characterize the happy moments in terms of affect, emotion, participants and content.

The task, given its predictive and open-ended interpretive aspects is relevant for the computational linguistics, natural language processing, artificial intelligence and the psycholinguistics communities. The aim is to engage scholarly interest and crowdsource new ideas and linguistic approaches to define happiness.

**Evaluation:** The performance of Systems was compared based on their Accuracy and F-1 measure at predicting the Agency and Sociality labels on the unseen test dataset. This was done using an automatic evaluation script, available on Github <sup>6</sup>.

## 1.1 Dataset description

The CL-Aff corpus comprises the following:

- **Labeled training set (N = 10,560):** Single-sentence happy moments from the available HappyDB corpus, annotated with demographic labels of the author, as well as labels that identify the 'agency' of the author and the 'social' characteristic of the moment, as well as concept labels describing its theme
- **Unlabeled training set (N = 72,324):** The remaining single-sentence HappyDB happy moments with only the demographic labels of the author.
- **Test set: (N = 17,327)** Previously unreleased, single-sentence happy moments, freshly collected in the same manner as the original HappyDB data. Authors' demographic labels were available to the Shared Task participants but not the 'agency' or 'social' characteristics.

The Agency and Sociality characteristics of each happy moment were decided by a simple majority agreement between three independent annotators using a binary (yes/no) coding.

## 2 Corpus development

### 2.1 Collecting the happy moments

We followed the format of the original HappyDB AMT task[1] to collect a second dataset of 20,000 happy moments, which was to be the unseen test data in the CL-Aff Shared Task. The following instructions were provided to the workers.

#### Instructions

What made you happy? Reflect on the past <duration>, and recall three actual events that happened to you that made you happy. Describe your happy moments with a complete sentence. Write three such moments. You will also

<sup>6</sup> <https://github.com/kj2013/claff-happydb/>

be asked to note for how long each event made you happy. This task also has post-task questions. Please be sure to answer the questions. Examples of happy moments we are NOT looking for (e.g., events in distant past, incomplete sentence): *The day I married my spouse*; *My dog*.

< Enter moment here >

For how long did that event make you happy? Select the answer that is most appropriate.

Each AMT worker was required to enter three happy moments experienced within a specific time period. Half of the questionnaires specified a time period of 24 hours, while the other half with a <time period> of 3 months. The options provided for the follow-up question about the duration (i.e., the length) of happiness were ‘All day, I’m still feeling it,’ ‘Half a day,’ ‘At least one hour,’ ‘A few minutes’ or ‘Not Applicable.’ After the participant answered these questions, demographic information was collected about their country, age, gender (‘Male’, ‘Female’, ‘Other’, ‘Not Applicable’), marital status (‘single’, ‘married’, ‘divorced’, ‘separated’, ‘widowed’ or ‘Not Applicable’), and whether or not they have children (‘yes’, ‘no’).

### 3 Annotation

Annotators were required to annotate each moment along two binary dimensions – Agency and Sociality.

**Instructions** Read the following happy moment. Choose any of the following that applies:

**Agency: Is the author in control? YES/NO**

Examples of sentences where the author is in control (Answer is YES):

- “*I ran on the treadmill for 20 minutes straight when I could barely do 5 minutes 3 months ago.*”
- “*Going out to a special birthday lunch for my great-grandmother in law’s birthday.*”

Examples of sentences where the author is not in control (Answer is NO):

- “*My youngest daughter got accepted to many prestigious universities and accepted an offer to attend college in San Diego.*”
- “*A small business deal change over for small profit.*”

**Social: Does this moment involve other people other than the author? YES/NO**

Please note that objects (e.g., bus, work) should not be counted as social. Examples of sentences which involve other people (Answer is YES):

- “*Going out to a special birthday lunch for my great-grandmother in law’s birthday.*”

- “My youngest daughter got accepted to many prestigious universities and accepted an offer to attend college in San Diego.”

Note that sometimes a person is implicitly involved although not explicitly mentioned. In this case, we still wish to label the happy moment as social. E.g., “I received compliments on my tattoo.”

Examples of sentences which are not social (Answer is NO):

- “I ran on the treadmill for 20 minutes straight when I could barely do 5 minutes 3 months ago.”
- “The bus came on time, so I reached work early.”

< *Happy moment appears here* >

**Agency: Is the author in control? YES/NO**

**Social: Does this moment involve other people other than the author? YES/NO**

### 3.1 Topic labeling

Annotators were presented with a happy moment, and a set of four potential topics which it was likely describing. Annotators were asked to mark all the tags which referred to what it was about. Each moment could score a maximum of four tags if at least two annotators agreed on them.

**Instructions** Read the following text. Select all categories that are relevant to the text from among those provided. If none of the categories is a great fit, select “none of the above”

<Topic 1> <Topic 2> <Topic 3> <Topic 4>

## 4 Overview of Approaches

Eleven teams participated in Task 1, with a subset of 2 of these additionally participating in Task 2. The following paragraphs discuss the approaches followed by the participating systems, in lexicographic order by team name.

- Team 1 [6]: The team from Arizona proposes a Word Pair Convolutional Model (WoPCoM) to accomplish Task 1. The proposed model is motivated behind a hypothesis that a small set of word-pair features are important to capturing the agency/social nature of happy moments. They trained a convolutional neural network (CNN) to predict on the unlabeled data.
- Team 6 [11]: The UC Santa Cruz team participated in both tasks. For Task 1, they explore using different feature sets to train their model, including syntactic features, emotional features, and survey features. They used semi-supervised learning and experimented with XGBoosted Forest and CNN

models. For Task 2, the team trained similar models to predict concepts, and based on the difficulty of doing so, hypothesized about the nature of the themes in the happy moments.

- Team 19 [8]: The IIIT Hyderabad team employs an inductive transfer learning technique (ITL). They pre-train a AWD-LSTM neural net on the WikiText-103 corpus, and then introduce an extra step to adapt the model to Happy moments.
- Team 21 [7]: The team from Gyrfalcon Technology, California, proposes an algorithm to map English words into squared glyphs images. Then, they apply a 2D-CNN model over these images in order to capture the sentiment.
- Team 22 [4]: IHPC-A\*STAR participated in both tasks. For Task 1, they use emotion intensity in happy moments to predict agency and sociality labels. They define a set of five emotions (valence, joy, anger, fear, sadness) and use a tool CrystalFeed that they developed in a previous work in order to label each moment with the corresponding five emotion intensities. Combining these features with additional word-embedding features, they train a logistic regression model. For Task 2, the team explores how these different emotions are manifested across the different concept labels.
- Team 24 [5]: The University of British Columbia team primarily experiments with using various deep neural nets, including long short-term memory networks and BiLSTM, with and without attention. They experiment using different embedding methods on these nets, such as CoVe and ELMo.
- Team 26 [12]: The University of Ottawa team also proposes a deep learning CNN solution. They experiment using different kind of word embeddings, and also experiment training a multi-task classifier to see if performance can be enhanced by shared knowledge between agency and sociality.
- Team 30 [10]: The Escuela Superior Politecnica del Litoral (ESPOL) team proposes a semi-supervised adaptation to traditional k-means clustering using neural networks.
- Team 31 [2]: The Sungkyunkwan team from South Korea uses a semi-supervised approach. They build four one-class autoencoder models, one for social, non-social, agentic, and non-agentic moments. Each autoencoder model has a deep learning architecture consisting of two neural networks, one for encoding the the input, and the other for reconstructing the compressed vector.
- Team 32 [9]: The Jordan University of Science and Technology team proposes their IoH-RCNN system. They use a Recurrent Convolutional Neural Network, and combine words with their context in order to get a more precise word embedding.
- Team 33 [3]: Finally, the team from Fraunhofer FKIE, Germany trained a three-layer CNN. They go through measures to clean out the corpus further, and experiments with using different embeddings including FastText and GloVe. Additionally, they experiment with splitting the dataset by demographic location of the author, and show that training separate classifiers on the splits enhances performance.

## 5 Results

### 5.1 Task 1: Predicting Agency and Sociality

This section compares the participating systems in terms of their performance. Four of the eleven systems that did Task 1 also did the bonus Task 2. The results are provided in Table 1. The detailed implementation of the individual runs are described in the system papers included in this proceedings volume.

### 5.2 Task 2: Happiness Insights

Some of the systems used their neural models of happiness for Task 1 to produce visual knowledge representations [7], and general insights about happiness [6,5,11,3]. Most notably, System 21 [7] transformed textual moments into visualizations to see if they could encode more multi-dimensional information in this manner. System 24 [5] provided a visualization for “attention” in their bi-directional long short-term memory networks which highlights the patterns that were considered important by the neural network while predicting Agency and Sociality when a sequence of words was input into the model. Team 1 [6] showed the codependence of the individual Agency and Sociality labels across the dataset through a t-SNE visualization. Team 33 [3] and Team 6 [11] both attempted to capture the linguistic patterns in the construction of happiness and their potential cultural underpinnings.

## 6 Error Analysis

In this section, we present a meta-analysis of system performances for Task 1 over all the (a) topics and (b) moments in the test set. Furthermore, in their data pre-processing step, Team 33 [3] identified that in the subset of happy moments contributed by authors from India alone, there were duplicate or near-duplicate happy moments in the data, which reduced the total number of training samples by 25%. We will include data cleaning as an extra preprocessing step in future data releases.

**Topic-level analysis:** We expect that happiness in different situations would be experienced and expressed differently. Table 3 aggregates the failures produced by each of the approaches (out of the set of best approaches submitted by each of the teams).

**Moment-level meta-analysis:** We suspect that some of the errors in our data may occur due to mislabeling or the coding scheme not being applicable to the moment. In Table 4 we provide the happy moments for which 100% of the best approaches submitted by each of the teams reported failure. We observe that in some of the cases, (e.g., “Topanga running away to Cory”), the happy moment was actually mislabeled and thus the systems actually did have the correct prediction. Overall, many of the happy moments in this Table describe a single moment in the author’s life, which seem ordinary when considered in the

Table 1: Systems’ performance in Task 1, ordered by their accuracy on predicting Agency.

System	Agency		Sociality	
	Accuracy	$F_1$	Accuracy	$F_1$
[5]	.85	.9	.92	.93
[6]	.85	.89	.91	.92
[8]	.84	.89	.92	.93
[9] Run 10	.83	.89	.91	.92
[11] Run 1	.83	.88	.89	.9
[11] Run 2	.83	.88	.89	.9
[11] Run 3	.83	.88	.89	.9
[9] Run 1	.83	.88	.91	.91
[9] Run 4	.83	.88	.91	.92
[9] Run 5	.83	.88	.9	.91
[9] Run 7	.82	.88	.91	.92
[9] Run 3	.82	.88	.9	.91
[9] Run 2	.82	.88	.91	.92
[10]	.82	.87	.9	.91
[9] Run 6	.82	.87	.9	.91
[11] Run 4	.82	.87	.89	.9
[12] Run 1	.82	.88	.9	.91
[4] Run 3	.81	.88	.89	.9
[12] Run 2	.81	.87	.89	.91
[9] Run 8	.81	.87	.89	.91
[12] Run 3	.81	.86	.88	.9
[12] Run 4	.8	.86	.88	.9
[9] Run 9	.8	.85	.9	.91
[12] Run 5	.8	.86	.89	.9
[12] Run 6	.8	.87	.89	.9
[12] Run 7	.8	.86	.88	.89
[12] Run 8	.8	.86	.89	.9
[7]	.8	.86	.82	.83
[12] Run 9	.79	.85	.88	.9
[11] Run 5	.79	.86	.59	.62
[12] Run 10	.79	.85	.88	.89
[11] Run 6	.79	.86	.7	.74
[4] Run 2	.78	.83	.89	.9
[4] Run 1	.78	.83	.89	.9
[3] Run 4 4	.77	.84	.65	.73
[3] Run 1	.76	.85	.59	.62
[11] Run 7	.76	.85	.89	.9
[11] Run 8	.76	.85	.89	.9
[3] Run 3 3	.76	.84	.65	.74
[3] Run 5	.76	.82	.62	.68
[2] <sup>7</sup> Run 1	.76	.84	.87	.88
[2] Run 2	.76	.84	.87	.88
[3] Run 2	.75	.82	.59	.62
[3] Run 6	.75	.84	.61	.65
[11] Run 9	.74	.83	.61	.66
[2] Run 4	.39	.49	.87	.88
[2] Run 3	.39	.49	.87	.88

Table 2: Legend for Task 1 System Runs.

System No.	Run No.	Description
[11]	Run 1	Supervised CNN (2,3,4,5) on GloVe features
[11]	Run 2	Semi-Supervised CNN (2,3,4) on GloVe features
[11]	Run 3	Semi-Supervised CNN (2,3,4,5) on GloVe features
[11]	Run 4	Supervised CNN (2,3,4) on GloVe features
[11]	Run 5	Supervised CNN (2,3,4) on GloVe, syntactic and emotionl features
[11]	Run 6	Semi-Supervised CNN (2,3,4,5) on GloVe features
[11]	Run 7	Supervised XGBoosted Forest on syntactic and emotion features
[11]	Run 8	Semi-Supervised SGBosted Forest on syntactic and emotion features
[11]	Run 9	Semi-Supervised CNN (2,3,4) on GloVe, syntactic and emotion features
[12]	CNN Multi-task learning on GloVe embeddings	
[12]	Run 1	50 dimensions, 10 epochs
[12]	Run 2	100 dimensions, 10 epochs
[12]	Run 3	25 dimensions, 10 epochs
[12]	Run 4	200 dimensions, 10 epochs
[12]	Run 5	25 dimensions, 50 epochs
[12]	Run 6	50 dimensions, 50 epochs
[12]	Run 7	100 dimensions, 50 epochs
[12]	Run 8	200 dimensions, 50 epochs
[12]	Run 9	50 dimensions, 200 epochs
[12]	Run 10	200 dimensions, 100 epochs

context of regular living. In some cases, the authors have attempted to explain why this moment was special to them (e.g., the second part of the moment “I finally got a hold of my auto mechanic, and that enabled me to schedule a time to bring in my car to get my custom exhaust installed” only serves to explain the significance of the moment to the author.)

## 7 Conclusion and Future Work

Eleven teams participated in Task 1 of the inaugural CL-Aff Shared Task AAAI-19 and out of those, five attempted Task 2. We have already published our complete dataset to Harvard Dataverse <sup>8</sup>. Furthermore, we expect to release other resources complementary to the challenges of modeling affect and emotion language from language.

We believe that the widespread adoption of neural approaches for modeling Agency and Sociality and the stupendous performance even on the modest size of the dataset, indicates the swift improvements happening in the field of deep learning for text. Our analysis identifies the following key takeaways:

- Predictive modeling approaches are greatly improved when modeled as a semi-supervised task, enriched with unlabeled data or by knowledge or feature vectors trained from a different domain. This also highlights the generalizability of the Shared Task to other domains.

<sup>8</sup> DOI:10.7910/DVN/JZAS66; <https://goo.gl/3rcZqf>



Table 3: Topic-level error Analysis: % of the approaches that failed on predicting Agency and Sociality in different topics.

Concept	N	% Failure (Agency)	% Failure (Sociality)
Career	2186	27.27	27.27
Party	839	27.27	36.36
Education	985	18.18	27.27
Family	5149	18.18	27.27
Animals	882	18.18	36.36
Religion	105	18.18	36.36
Conversation	1336	18.18	27.27
Romance	1325	18.18	27.27
Weather	251	18.18	18.18
Vacation	1061	9.09	27.27
Entertainment	2284	9.09	27.27
Food	2402	9.09	18.18
Shopping	1290	9.09	18.18
Technology	477	9.09	18.18
Exercise	842	9.09	18.18

Table 4: Moment-level error Analysis: 100% of the best approaches submitted by each of the eleven teams, failed at predicting Agency and Sociality in these happy moments.

Agency (10% noise)
<p>I was given off phone work for the day.</p> <p>i was promoted in my job and i felt so appreciated and happy.</p> <p>Children of different races playing together at pool.</p> <p>I spoke to a friend on the phone that I hadn't heard from in two months.</p> <p>I realized earlier today that last month I made twice as much as my income for the same month the previous year.</p> <p>I won a new lawnmower, which I desperately needed, at a local event.</p> <p>My son singing songs with me.</p> <p>seeing a co worker</p> <p>I was able to talk to my boyfriend on the phone for an hour even though he is on vacation.</p> <p>i slept well last night - no nightmares</p> <p>I got a raise at work</p> <p>My father gifted me a car on my birthday that left me surprised and extremely happy.</p> <p>I won a lucky draw two weeks back to a five star resort for two days, which made me really exiting and joyfull.</p>
Social (5% noise)
<p>I was applying for jobs for many months and finally got an interview and a offer later.</p> <p>I learned that we're moving into a bigger, better apartment for less money.</p> <p>I won the first prize in cricket match.</p> <p>I was happy when I was offered to work on a new television show</p> <p>Having my hair cut and it turning out just the way I wanted it to look.</p> <p>I had a job interview this morning and I think it went well.</p> <p>Topanga running away to Cory.</p> <p>Being brought McDonald's for lunch.</p> <p>I was gifted a very nice bottle of wine.</p> <p>I got an iphone 7 as gift</p> <p>I forgot my credit card and was given a free chicken biscuit and drink anyway.</p> <p>A package I ordered that I was anticipating was delivered to me.</p> <p>I received approval to take a month off of work to go on a backpacking trip.</p> <p>Got a small raise at work and even it doesn't amount to much per check, it's still something.</p> <p>When i return home my house was very clean and nice,and the yard was mowed.</p>

- Many of the teams were able to identify and model the importance of syntactic knowledge for predicting Agency and Sociality. They did so by recognizing the importance of the head noun and subject-verb-object word order, either by modeling it through interacting layers in convolution neural networks, or by mining it using lexical pattern analyses methods.
- One of the teams was able to apply the models developed on other annotated emotion datasets to improve their predictive modeling performance on the Shared Task dataset [4]. It is encouraging to find out that the dataset offers replicability of more traditional emotion modeling approaches.

The participating teams suggested that predictive modeling would be complemented with domain-specific word embeddings trained on the Shared Task dataset itself. Many of the teams successfully used task-specific embeddings from other language models, such as ELMo and CoVe, which encourages the notion of thinking about words in the English language in a contextualized emotional vector space. We anticipate that language models from the CL-Aff dataset will also generalize well to other problems and datasets for emotion and affect analysis.

**Acknowledgement.** We thank Wang-Chiew Tan for her feedback and Megagon Labs for contributing funds towards the CL-Aff dataset.

## References

1. Asai, A., Evensen, S., Golshan, B., Halevy, A., Li, V., Lopatenko, A., Stepanov, D., Suhara, Y., Tan, W.C., Xu, Y.: Happydb: A corpus of 100,000 crowdsourced happy moments. In: Proceedings of LREC 2018. European Language Resources Association (ELRA), Miyazaki, Japan (May 2018)
2. Cheong, Y.G., Song, Y., Bae, B.C.: [cl-aff shared task] modeling happiness using one-class autoencoders. In: Proceedings of the 2nd Workshop on Affective Content Analysis @ AAAI (AffCon2019). Honolulu, Hawaii (January 2019)
3. Claeser, D.: Affective content classification using convolutional neural networks. In: Proceedings of the 2nd Workshop on Affective Content Analysis @ AAAI (AffCon2019). Honolulu, Hawaii (January 2019)
4. Gupta, R.K., Bhattacharya, P., Yang, Y.: What constitutes happiness? predicting and characterizing the ingredients of happiness using emotion intensity analysis. In: Proceedings of the 2nd Workshop on Affective Content Analysis @ AAAI (AffCon2019). Honolulu, Hawaii (January 2019)
5. Rajendran, A., Zhang, C., Abdul-Mageed, M.: Happy together: Learning and understanding appraisal from natural language. In: Proceedings of the 2nd Workshop on Affective Content Analysis @ AAAI (AffCon2019). Honolulu, Hawaii (January 2019)
6. Saxon, M., Bhandari, S., Ruskin, L., Honda, G.: Word pair convolutional model for happy moment classification. In: Proceedings of the 2nd Workshop on Affective Content Analysis @ AAAI (AffCon2019). Honolulu, Hawaii (January 2019)
7. Sun, B., Yang, L., Chi, C., Zhang, W., Lin, M.: [cl-aff shared task] squared english word: A method of generating glyph to use super characters for sentiment analysis. In: Proceedings of the 2nd Workshop on Affective Content Analysis @ AAAI (AffCon2019). Honolulu, Hawaii (January 2019)

8. Syed, B., Indurthi, V., Shah, K., Gupta, M., Varma, V.: Ingredients for happiness: Modeling constructs via semi-supervised content driven inductive transfer learning. In: Proceedings of the 2nd Workshop on Affective Content Analysis @ AAAI (AffCon2019). Honolulu, Hawaii (January 2019)
9. Talafha, B., Al-Ayyoub, M.: Ioh-rcnn: Pursuing the ingredients of happiness using recurrent convolutional neural networks. In: Proceedings of the 2nd Workshop on Affective Content Analysis @ AAAI (AffCon2019). Honolulu, Hawaii (January 2019)
10. Torres, J., Vaca, C.: Neural semi-supervised learning for short-texts. In: Proceedings of the 2nd Workshop on Affective Content Analysis @ AAAI (AffCon2019). Honolulu, Hawaii (January 2019)
11. Wu, J., Compton, R., Rakshit, G., Walker, M., Anand, P., Whittaker, S.: Cruzafect at affcon 2019 shared task: A feature-rich approach to characterize happiness. In: Proceedings of the 2nd Workshop on Affective Content Analysis @ AAAI (AffCon2019). Honolulu, Hawaii (January 2019)
12. Xin, W., Inkpen, D.: [cl-aff shared task] happiness ingredients detection using multi-task deep learning. In: Proceedings of the 2nd Workshop on Affective Content Analysis @ AAAI (AffCon2019). Honolulu, Hawaii (January 2019)